



UNITED STATES INTERNATIONAL UNIVERSITY – AFRICA
(USIU-Africa)

School of Science and Technology
Department of Computing



END OF SEMESTER PROJECT REPORT

PROJECT

Smart Health Symptom Checker and Doctor Recommendation System

Course	SWE3090XA – Software Engineering Final Project
Submitted by	Amy Wanjugu Njau
Registration number	669008
Supervisor	Mr Fredrick Ogore
Submission date	August 2026

Summer Semester 2026

Table of Contents

Abstract	4
Chapter One: Introduction	5
1.1 Background of the Study	5
1.2 Emerging Issues	5
1.3 Problem Statement	5
1.4 Objectives	5
1.5 Justification and Significance	6
1.6 Scope and Limitations	6
Chapter Two: Literature Review	7
2.1 Review of Existing Systems	7
2.2 Comparative Analysis	7
2.3 Identified Gaps	7
2.4 Unique Contributions	8
Chapter Three: Methodology	9
3.1 Development Approach	9
3.2 Project Timeline	9
3.3 Resources and Tools	9
3.4 Requirements	9
Chapter Four: System Design	11
4.1 Architecture	11
4.2 Context Diagram	11
4.3 Data Flow Diagram	12
4.4 Entity Relationship Diagram	12
4.5 API Contract	13
4.6 Security and Privacy Design	14
Chapter Five: Implementation	15
5.1 Codebase Structure	15
5.2 Symptom-Condition Knowledge Base	15
5.3 Diagnostic Engine	15
5.4 API Endpoint	16
5.5 Provider Recommendation	16
5.6 Mobile Client	16
5.7 User Interface	17
5.8 Key Modules and Algorithms	19
Chapter Six: Testing and Validation	20
6.1 Testing Strategy	20
6.2 Automated Test Results	20
6.3 Acceptance Test Cases	20
6.4 Non-Functional Verification	21
Chapter Seven: Results and Discussion	22
7.1 Objectives Achieved	22
7.2 Discussion	22

7.3 Limitations	22
Chapter Eight: Conclusions	23
8.1 Conclusions	23
8.2 Recommendations	23
8.3 Closing Remarks	23
References	24
Appendices	25
Appendix A: Repository	25
Appendix B: Running the System	25
Appendix C: Knowledge Base Coverage	25

Abstract

Access to timely medical guidance remains constrained in many developing regions, where low doctor-to-patient ratios, travel distance, and consultation cost combine to delay care. This project delivers the Smart Health Symptom Checker and Doctor Recommendation System: a cross-platform mobile application that accepts user-reported symptoms, returns a ranked list of probable conditions together with the reasoning behind each ranking, recommends the appropriate medical specialist, and locates nearby qualified providers using the user's real-time position.

The system is built on a three-tier architecture: a Flutter mobile client, a Node.js and Express REST API hosting a rule-based diagnostic engine, and a swappable data and location tier. The diagnostic engine scores each candidate condition by summing the weights of the symptoms the user reported and dividing by that condition's maximum possible weight, producing a normalised confidence value. Every score therefore traces back to specific named symptoms, which the interface surfaces to the user. This explainability is the project's central design commitment and the principal reason a rule-based approach was chosen over an opaque machine-learning model in a safety-sensitive domain.

The delivered system implements a knowledge base of 15 conditions, 35 symptoms, and 9 specialist types, exposes five REST endpoints, and is verified by an automated suite of 15 tests, all of which pass. Functional and non-functional requirements defined at the midterm stage were met, including deterministic results, graceful degradation when location permission is denied, and a guidance disclaimer attached centrally to every diagnosis response so that it cannot be omitted. The report documents the requirements, design, implementation, testing, and results, and sets out the security hardening, clinical validation, and persistence work required before any real-world deployment.

Chapter One: Introduction

1.1 Background of the Study

Access to timely and accurate medical guidance remains a major challenge across many parts of the world, and particularly in developing regions where the ratio of doctors to patients is critically low. Large populations in low- and middle-income countries continue to face shortages of healthcare workers, long travel distances to facilities, and high out-of-pocket consultation costs. As a result, many individuals delay seeking professional medical attention, not because they are unwilling, but because they are uncertain about the severity of their symptoms or do not know which type of practitioner to consult.

At the same time, mobile-phone penetration in these regions has grown rapidly, placing internet-connected devices in the hands of populations that were previously difficult to reach through conventional healthcare infrastructure. This convergence creates a meaningful opportunity for mobile health solutions to bridge part of the gap between symptom onset and appropriate care.

The Smart Health Symptom Checker and Doctor Recommendation System responds to this opportunity by providing a digital-first tool that helps users make informed health decisions. By entering their symptoms into a mobile application, users receive condition suggestions generated through a transparent rule-based analysis engine. The system then advises which type of medical specialist to consult, and recommends nearby qualified providers based on the user's real-time location. The application acts as an accessible first point of guidance: not a replacement for professional care, but a bridge toward it.

1.2 Emerging Issues

Several emerging issues motivate and shape this project.

Rising demand for self-directed care. Patients increasingly turn to online sources to interpret their symptoms before deciding whether and where to seek help. Unstructured web searches, however, often produce alarming or contradictory information.

Proliferation of digital symptom checkers. A growing number of commercial symptom-checking applications exist, but many are designed for high-income contexts, require constant connectivity, or do not connect the user to local, reachable care.

Localisation gap. Few existing tools account for the realities of developing regions: nearby facilities, locally available specialists, and low-bandwidth conditions.

Trust, safety, and explainability. As automated health tools become more common, there is increased scrutiny of how recommendations are generated. Rule-based systems offer transparency that opaque models cannot easily match, which is valuable in a sensitive domain such as health.

Data privacy expectations. Handling even limited health-related data raises legitimate concerns around consent, storage, and confidentiality that any responsible solution must address.

1.3 Problem Statement

Many individuals lack access to timely, affordable, and location-aware medical guidance, leading to delayed diagnoses and preventable deterioration of their health. Existing solutions are either too generic to be trustworthy, too expensive for the average user, or not localised to the user's context, failing to connect them to qualified care that is actually within reach.

There is therefore a need for an accessible mobile application that can interpret a user's symptoms, suggest probable conditions in an explainable manner, recommend the appropriate type of specialist, and direct the user to nearby qualified healthcare providers, all through a simple interface usable by people with varying levels of digital literacy.

1.4 Objectives

1.4.1 MAIN OBJECTIVE

To design and develop a cross-platform mobile application that accepts user-reported symptoms, returns a ranked list of probable medical conditions, recommends an appropriate medical specialist, and identifies nearby qualified healthcare providers based on the user's real-time location.

1.4.2 SPECIFIC OBJECTIVES

1. To design and develop a mobile application that accepts user-inputted symptoms and returns a ranked list of possible medical conditions.

2. To implement a rule-based diagnostic engine capable of mapping symptom combinations to probable diseases with reasonable accuracy.
3. To integrate a specialist recommendation module that suggests the appropriate type of doctor based on the predicted condition.
4. To incorporate location-based services that identify and recommend nearby qualified healthcare providers or medical facilities.
5. To provide a simple, intuitive user interface accessible to users with varying levels of digital literacy.

1.5 Justification and Significance

This project is justified by the persistent gap between the demand for healthcare guidance and the supply of accessible professional care in developing regions. By combining symptom analysis with localised, location-aware provider recommendation, the system addresses a need that existing tools only partially serve.

For users, it provides a free, easy-to-use first point of guidance that reduces uncertainty, encourages earlier care-seeking, and helps users reach the right kind of practitioner faster. For the healthcare system, guiding patients toward the appropriate specialist and nearby facilities can reduce mis-directed visits and ease pressure on overstretched general services. For underserved communities, the emphasis on localisation and a lightweight, explainable approach suits low-connectivity environments that commercial alternatives often overlook. Academically, the project demonstrates the practical application of software-engineering principles, including requirements engineering, system design, rule-based inference, API integration, automated testing, and Agile delivery, to a real-world problem of social value.

1.6 Scope and Limitations

The delivered system covers the full path from symptom entry to nearby-provider recommendation on Android, with an architecture that permits an iOS build without a rewrite. It does not extend to telemedicine consultation, appointment booking, electronic health records, payments, or prescription handling.

Three limitations are stated plainly here and revisited in Section 7.3. First, the knowledge base is an educational rule set curated for this project and has not been clinically validated; the application is a guidance tool and not a medical device. Second, the query log is held in memory and is therefore lost when the server restarts. Third, the deployed prototype has no user authentication, so history is device-local rather than per-account.

Chapter Two: Literature Review

This chapter reviews existing symptom-checker and digital-health systems, compares their capabilities against the goals of this project, identifies the gaps they leave open, and states the unique contributions of the delivered system.

2.1 Review of Existing Systems

2.1.1 WEBMD SYMPTOM CHECKER

WebMD provides a widely used web-based symptom checker that allows users to select symptoms on a body map and returns a list of possible conditions with accompanying articles. It is information-rich but primarily an educational reference; it does not connect users to specific local practitioners and is oriented toward a North American audience and healthcare context.

2.1.2 ADA HEALTH

Ada is a mobile symptom-assessment application that conducts a structured, conversational interview and returns possible causes together with guidance on next steps. It is well designed and clinically informed, but it is built primarily for connected, higher-income markets, does not focus on surfacing nearby local providers, and its underlying assessment logic is largely opaque to the end user.

2.1.3 BABYLON HEALTH

Babylon combined symptom checking with telemedicine consultations, allowing users to speak with clinicians remotely. While powerful, it depends on reliable connectivity, an available clinician network, and a subscription or payment model, which limits its accessibility in low-resource settings.

2.1.4 BUOY HEALTH AND K HEALTH

Buoy Health and K Health both use data-driven approaches to triage symptoms and suggest next steps, with K Health additionally drawing on large clinical datasets. These tools illustrate the value of structured triage but, like the others, are tailored to specific national healthcare systems and do not provide localised facility recommendation for developing-region users.

2.1.5 ACADEMIC EVALUATIONS

Independent academic audits of symptom checkers, notably the BMJ audit study by Semigran and colleagues (2015), found that the diagnostic and triage accuracy of these tools varies considerably. That body of work underpins two design principles adopted here: such tools should be positioned explicitly as guidance rather than diagnosis, and an explainable, well-validated rule base is preferable to an opaque one in a safety-sensitive domain.

2.2 Comparative Analysis

The table below compares representative existing systems against the core capabilities targeted by this project.

System	Symptom analysis	Specialist recommendation	Location-aware providers	Explainable logic	Suited to low connectivity
WebMD Symptom Checker	Yes	Partial	No	Partial	No
Ada Health	Yes	Partial	No	No	No
Babylon Health	Yes	Yes, via clinician	Partial	No	No
Buoy Health / K Health	Yes	Partial	No	No	No
Smart Health (this project)	Yes	Yes	Yes	Yes	Yes

Table 2.1: Comparative analysis of existing symptom-checker systems.

2.3 Identified Gaps

The review reveals four consistent gaps. Most tools are designed for specific high-income healthcare systems and do not surface nearby, reachable providers relevant to a developing-region user. Few systems explicitly translate a probable condition into a clear recommendation of which type of specialist to consult. Many modern tools provide little visibility into how a

recommendation was produced, which undermines trust in a sensitive domain. Finally, solutions depending on continuous connectivity, large models, or paid consultations are poorly suited to low-resource environments.

2.4 Unique Contributions

The delivered system addresses those gaps through four contributions. It provides integrated end-to-end guidance in a single workflow, moving the user from symptom entry to probable condition, to the right specialist, to a nearby provider. It uses real-time geolocation to recommend providers within reach of the user. It applies explainable rule-based inference over a transparent weighted symptom-condition matrix whose reasoning is surfaced in the interface rather than hidden. And it is deliberately lightweight, targeting varying digital-literacy levels and lower-connectivity conditions.

Chapter Three: Methodology

3.1 Development Approach

The project followed an Agile software-development methodology implemented through short, iterative sprints. Agile was selected over a strictly sequential approach such as Waterfall because a health-guidance tool benefits from continuous testing, frequent feedback, and incremental refinement, particularly for the diagnostic rule base, which had to be validated and tuned progressively rather than fixed up front.

Development was organised into four phases. The requirements and research phase reviewed existing symptom-checker solutions, defined the symptom-condition rule base, and settled the architecture. The design phase produced the data models, the API contract, and the inference design. The development and integration phase built the mobile application, implemented the diagnostic engine, added geolocation, and connected the backend. The testing phase covered unit, integration, and acceptance testing.

3.2 Project Timeline

The schedule spanned approximately ten weeks. All milestones are complete.

Week	Activity	Deliverable	Status
1-2	Requirements gathering and literature review	Requirements document	Complete
3-4	System design, data models, API contract	Design models and wireframes	Complete
5-6	Database design, backend setup, rule-engine development	Working rule engine	Complete
7-8	Mobile UI development: symptom input and results	Alpha build	Complete
9	Geolocation integration and specialist recommendation	Beta build with location	Complete
10	End-to-end testing, fixes, documentation, submission	Tested application and report	Complete

Table 3.1: Project timeline and final status.

3.3 Resources and Tools

Development used a personal computer, an Android device, and the Android emulator for testing. The software stack is summarised below.

Category	Technology	Purpose
Frontend	Flutter (Dart)	Cross-platform mobile interface, Android first, iOS-ready
Backend	Node.js with Express	REST API and integration mediator
Diagnostic logic	Rule-based inference engine	Weighted symptom-condition scoring
Data	JSON data store, Firebase Firestore ready	Knowledge base, specialists, providers, query log
Location	Google Maps Platform, device geolocation	Nearby provider discovery
Testing	Node.js built-in test runner	Automated unit and integration tests
Version control	Git and GitHub	Source-code management
Tooling	VS Code, Android Studio	Development, debugging, emulation

Table 3.2: Tools and technologies.

3.4 Requirements

3.4.1 FUNCTIONAL REQUIREMENTS

The system shall allow users to input symptoms through structured selection; analyse them with the rule-based engine and return a ranked list of probable conditions; recommend the appropriate specialist for the top-ranked condition; detect the user's location with permission and recommend nearby providers; display a clear disclaimer that the application provides guidance

only; and handle denial of location permission gracefully while still providing symptom and specialist guidance.

3.4.2 NON-FUNCTIONAL REQUIREMENTS

Usability requires an interface that is simple for users with varying digital literacy. Performance requires symptom analysis to return within a few seconds on a typical mobile connection. Reliability requires the engine to produce consistent results for identical input and to degrade gracefully when services are unavailable. Security and privacy require health and location data to be handled with consent, minimisation, and confidentiality. Portability requires an Android-first architecture that permits iOS extension without a rewrite. Maintainability requires the rule base to accept new symptoms, conditions, and specialists without changes to core logic.

Chapter Four: System Design

4.1 Architecture

The system adopts a three-tier architecture. The presentation layer is the Flutter mobile application, responsible for symptom input, displaying ranked results and specialist advice, and listing nearby providers. The application layer is the Node.js and Express backend, which exposes REST endpoints, hosts the rule-based engine, performs specialist mapping, and orchestrates calls to the location service. The data layer holds the knowledge base, specialist mappings, provider records, and an anonymised query log, with the Google Maps Platform supplying external place data.

A deliberate design decision is that the mobile client contains no diagnostic logic whatsoever. It collects input, calls the API, and maps JSON responses into typed models. All reasoning lives behind the API boundary, which keeps the client thin, makes the engine independently testable, and allows the rule base to be corrected without shipping a new app version.




Figure 4.1: High-level system architecture.

4.2 Context Diagram

The context diagram represents the system as a single process interacting with its external entities: the user, who submits symptoms and location permission and receives guidance; the Google Maps Platform, which receives location and specialty queries and returns nearby places; and the data store, which supplies symptom-condition rules and specialist mappings.




Figure 4.2: Context diagram, level-0 data flow.

4.3 Data Flow Diagram

The level-1 diagram decomposes the system into five processes: capture symptoms, analyse symptoms with the rule engine, recommend a specialist, recommend nearby providers, and present results. Data stores comprise the knowledge base, the specialist mapping, the provider store, and the anonymised query log.




Figure 4.3: Level-1 data flow diagram.

4.4 Entity Relationship Diagram

The data model is organised around seven entities. The associative entity Condition_Symptom carries the rule weight, which is what makes the knowledge base tunable without schema changes.

Entity	Key fields	Relationship
User	user_id (PK), name, location, created_at	Initiates many queries
Symptom	symptom_id (PK), name, category	Linked to conditions, many-to-many
Condition	condition_id (PK), name, description, severity	Linked to symptoms; maps to one specialist
Condition_Symptom	condition_id (FK), symptom_id (FK), weight	Associative entity holding the rule weight
Specialist	specialist_id (PK), type, description	Recommended by many conditions
Provider	provider_id (PK), name, specialty, latitude, longitude, rating	Offers one or more specialist types
Query	query_id (PK), user_id (FK), symptoms, result, timestamp	Belongs to one user

Table 4.1: Core entities of the data model.




Figure 4.4: Entity relationship diagram.

4.5 API Contract

The API boundary is the integration contract between client and server. Five endpoints are exposed.

Method and path	Purpose
POST /api/diagnose	Accepts a symptom set; returns ranked conditions, specialist, disclaimer
POST /api/providers	Accepts specialist and coordinates; returns ranked nearby providers
GET /api/symptoms	Returns the symptom catalogue for the input screen
GET /api/specialists	Returns the specialist catalogue
GET /api/health	Liveness check used by the test suite and monitoring

Table 4.2: REST API surface.

The end-to-end exchange is shown in Figure 4.5.




Figure 4.5: End-to-end data flow between client, API, engine and services.

4.6 Security and Privacy Design

Four measures were designed in rather than added afterwards. All client-server traffic is intended to travel over HTTPS. The Google Maps Platform key is held server-side in environment variables and is never exposed to the client, which is why every Maps call is made by the backend rather than the app. Request bodies are validated before processing, and malformed requests are rejected with a 400 status and a clear message. The error handler never leaks internal detail on a server error.

Data minimisation is applied to the query log: it records the symptom set and the outcome for analytics but is not tied to an identified person in the prototype. Location is read only after explicit user permission, and denial is a supported path rather than an error.

Section 7.3 records honestly the hardening that remains outstanding, since the prototype was assessed against a production security bar and did not yet meet it.

Chapter Five: Implementation

5.1 Codebase Structure

The repository is organised into three top-level areas: `backend/` for the API and engine, `mobile/` for the Flutter client, and `docs/` for the design figures and reports. Within the backend, the engine is isolated from all input and output so that it can be tested directly.

Path	Responsibility
<code>src/services/diagnosisEngine.js</code>	Pure scoring logic, no input or output
<code>src/services/diagnosisService.js</code>	Orchestrates engine, specialist lookup, query log, disclaimer
<code>src/services/providerService.js</code>	Mock and Google provider lookup with a cache
<code>src/routes/</code>	The five HTTP endpoints
<code>src/middleware/validate.js</code>	Request-body validation returning 400 responses
<code>src/repositories/</code>	Data-store factory, local and Firestore implementations
<code>src/data/</code>	Knowledge base, symptoms, specialists, providers
<code>src/config/index.js</code>	All environment reads and the central disclaimer

Table 5.1: Backend codebase structure.

5.2 Symptom-Condition Knowledge Base

The knowledge base encodes each condition as a set of weighted symptoms, where the weight expresses how strongly a symptom indicates that condition. Keeping it declarative means new conditions can be added without touching engine code, which satisfies the maintainability requirement directly.

```
{
  "id": "malaria",
  "name": "Malaria",
  "specialist": "General Physician",
  "severity": "high",
  "description": "A mosquito-borne infectious disease causing cyclical fever, chills, and flu-like illness.",
  "symptoms": {
    "fever": 0.9,
    "chills": 0.8,
    "headache": 0.5,
    "fatigue": 0.4,
    "nausea": 0.4,
    "sweating": 0.6,
    "muscle_aches": 0.4
  }
}
```

The delivered base holds 15 conditions spanning three severity levels, 35 symptoms, and 9 specialist types. Its metadata block states explicitly that the rule set is educational and not clinically validated.

5.3 Diagnostic Engine

The engine is the core contribution. For each condition it sums the weights of the symptoms the user reported, divides by that condition's maximum possible weight, and returns both the normalised score and the list of symptoms that contributed to it.

```
function scoreCondition(condition, userSymptoms) {
  const weights = condition.symptoms || {};
  const maxScore = Object.values(weights).reduce((a, b) => a + b, 0);
  if (maxScore <= 0) return { score: 0, matched: [] };

  let matchedWeight = 0;
  const matched = [];
  for (const symptom of userSymptoms) {
    if (Object.prototype.hasOwnProperty.call(weights, symptom)) {
      matchedWeight += weights[symptom];
      matched.push(symptom);
    }
  }
  return { score: matchedWeight / maxScore, matched };
}
```

Ranking then filters below a configurable threshold of 0.2, sorts by descending confidence, and caps the result count. Input is de-duplicated defensively before scoring, which is what makes repeated symptom submissions safe.

A worked example makes the arithmetic concrete. A user reporting fever, headache, and chills matches malaria on weights 0.9, 0.5, and 0.8, summing to 2.2. The maximum possible weight for malaria is 4.0. The normalised confidence is therefore $2.2 / 4.0$, or 55 per cent, and the interface reports exactly those three symptoms as the reason. Nothing about the number is hidden from the user.

5.4 API Endpoint

The route layer is deliberately thin: validation runs as middleware, the service performs the work, and errors are delegated to the central handler.

```
router.post('/', validateDiagnoseBody, async (req, res, next) => {
  try {
    const { symptoms } = req.body;
    const payload = await runDiagnosis(symptoms);
    res.json(payload);
  } catch (err) {
    next(err);
  }
});
```

5.5 Provider Recommendation

The provider service supports two interchangeable sources selected by configuration. The mock source ranks the bundled provider records by haversine distance from the user. The Google source queries the Maps Platform Nearby Search server-side using the specialist's search keywords, and caches responses for a configured period to contain both latency and billing. Because both sources satisfy the same interface, neither the routes nor the client change when the source is switched.

5.6 Mobile Client

The Flutter client presents a tabbed shell with Home, Check, History, and Profile. The symptom-input screen collects a symptom set, the results screen renders ranked conditions with their explainability chips, and the providers screen lists nearby practices with a working directions action. A single service class is the only point of contact with the API.

```
Future<DiagnosisResult> getDiagnosis(List<String> symptoms) async {
  final response = await http.post(
    Uri.parse('$baseUrl/api/diagnose'),
    headers: {'Content-Type': 'application/json'},
    body: jsonEncode({'symptoms': symptoms}),
  );
  if (response.statusCode == 200) {
    return DiagnosisResult.fromJson(jsonDecode(response.body));
  } else {
    throw Exception('Diagnosis request failed: ${response.statusCode}');
  }
}
```


5.7 User Interface




Figure 5.1: Home dashboard.




Figure 5.2: Symptom input screen.




Figure 5.3: Ranked conditions with the symptoms that produced each score.




Figure 5.4: Nearby provider recommendations.




Figure 5.5: Query history.

5.8 Key Modules and Algorithms

The knowledge-base module stores conditions declaratively so the rule set can grow without code changes. The engine module implements the weighted scoring described in Section 5.3 and is pure, which is what allows it to be unit-tested without a server. The specialist module translates the top-ranked condition into a concrete practitioner type, turning an abstract result into an actionable next step. The provider module converts that specialist type plus coordinates into a ranked shortlist of reachable facilities. The safety module attaches the guidance disclaimer centrally in configuration, so no response path can omit it, and severity is carried through to the interface so that high-severity results can be visually emphasised.

Chapter Six: Testing and Validation

6.1 Testing Strategy

Validation combined three levels. Unit tests exercise the diagnostic engine in isolation, confirming that scores are computed, normalised, filtered, and ranked correctly for known symptom sets. Integration tests drive the HTTP API with the engine, data store, and provider service wired together. Acceptance testing walked the complete user journey on an Android emulator against a live backend.

The engine was written as a pure function specifically so that the most safety-critical logic could be tested exhaustively without network, database, or user interface in the way.

6.2 Automated Test Results

The automated suite comprises 15 tests and is executed with `npm test`. All 15 pass. The suite is listed in full below because it constitutes the primary evidence that the functional and reliability requirements are met.

Test	Level	Requirement verified	Result
<code>GET /api/health</code> reports ok	Integration	Service liveness	Pass
<code>POST /api/diagnose</code> returns ranked conditions, specialist, disclaimer	Integration	Functional requirements 2, 3, 5	Pass
<code>POST /api/diagnose</code> rejects empty symptoms with 400	Integration	Input validation	Pass
<code>POST /api/providers</code> returns providers ranked by distance	Integration	Functional requirement 4	Pass
<code>POST /api/providers</code> rejects missing coordinates with 400	Integration	Input validation	Pass
<code>GET /api/symptoms</code> returns the catalogue	Integration	Symptom input support	Pass
<code>scoreCondition</code> returns 1.0 when all symptoms match	Unit	Scoring upper bound	Pass
<code>scoreCondition</code> normalises a partial match	Unit	Normalisation	Pass
<code>scoreCondition</code> ignores unknown symptoms	Unit	Robustness to bad input	Pass
<code>diagnose</code> ranks conditions by confidence descending	Unit	Ranking	Pass
<code>diagnose</code> drops conditions below the threshold	Unit	Threshold filtering	Pass
<code>diagnose</code> returns explainable matched symptoms	Unit	Explainability	Pass
<code>diagnose</code> is deterministic for the same input	Unit	Reliability requirement	Pass
<code>diagnose</code> de-duplicates repeated symptoms	Unit	Input hygiene	Pass
<code>diagnose</code> respects maxResults	Unit	Result capping	Pass

Table 6.1: Automated test suite, 15 of 15 passing.

6.3 Acceptance Test Cases

ID	Scenario	Expected outcome	Result
T1	Submit a valid symptom set	Ranked conditions and specialist returned	Pass
T2	Submit an empty symptom list	400 with a clear validation message	Pass
T3	Request providers with valid coordinates	Ranked nearby providers returned	Pass
T4	Deny location permission	Flow continues with symptom and specialist guidance	Pass
T5	Simulate network failure on the client	Readable error with retry, no crash	Pass
T6	Complete journey from symptoms to providers	Workflow completes without data loss	Pass

Table 6.2: Acceptance test cases.

6.4 Non-Functional Verification

Reliability is verified directly by the determinism test, which asserts that identical input produces identical output. Performance was observed on the emulator against a local backend, where diagnosis responses returned well within the few-seconds target; the engine itself completes in under a millisecond, as the unit-test timings show, so latency is dominated by network and rendering rather than by inference. Graceful degradation is verified by test T4, in which denying location permission still yields symptom and specialist guidance rather than an error state. Maintainability was demonstrated in practice: conditions were added to the knowledge base during development without a single change to engine code.

Chapter Seven: Results and Discussion

7.1 Objectives Achieved

Objective	Outcome	Evidence
Mobile app returning ranked conditions	Achieved	Figures 5.2 and 5.3; integration tests
Rule-based diagnostic engine	Achieved	Section 5.3; nine unit tests
Specialist recommendation module	Achieved	Diagnose response; Table 6.1
Location-based provider recommendation	Achieved	Figure 5.4; provider tests
Simple, accessible interface	Achieved	Figures 5.1 to 5.5; acceptance tests

Table 7.1: Achievement of the specific objectives.

All five specific objectives, and therefore the main objective, were met by the delivered system.

7.2 Discussion

The most important result is that explainability cost nothing in capability. The weighted rule base produced sensible rankings across the tested symptom sets while remaining fully inspectable: every percentage the user sees can be recomputed by hand from the knowledge base, as demonstrated in Section 5.3. In a domain where an unexplained recommendation may be ignored or, worse, trusted blindly, that property is worth more than a marginal accuracy gain from an opaque model.

The second result concerns architecture. Placing all reasoning behind an API boundary, and defining both the data store and the provider source as interchangeable implementations of a fixed interface, meant that the mock and live location paths could be switched by configuration alone. This is why the project could be developed and demonstrated reliably without incurring Maps billing, while remaining ready for real keys.

The third observation is about safety engineering. Attaching the disclaimer in central configuration rather than at each call site converts a discipline problem into a structural guarantee: no future endpoint can forget it. The same reasoning motivated making the engine pure, since the component with the highest consequence of failure became the easiest to test.

7.3 Limitations

Four limitations are stated plainly. The knowledge base is educational and has not been reviewed by a qualified clinician, so the system must not be used for real medical decisions. The query log is in memory and does not survive a restart, so the Firestore implementation remains a documented stub. The prototype has no authentication, so history is device-local. And the API, assessed against a production bar, still needs security hardening: origin restriction, security headers, rate limiting on the unauthenticated endpoints, bounds on request payloads, and eviction policies for the in-memory cache and log. These are enumerated as concrete work items rather than left implicit.

Chapter Eight: Conclusions

8.1 Conclusions

The project set out to close a specific gap: the absence of a tool that combines explainable symptom analysis, specialist routing, and localised provider recommendation in a form suited to low-connectivity, developing-region contexts. The delivered system does so end to end, and its behaviour is backed by an automated suite of 15 tests, all passing.

The rule-based approach proved to be the correct choice for this problem. It produced ranked, defensible results; it made the reasoning available to the user instead of hiding it; and it kept the highest-risk component small enough to test exhaustively. The three-tier architecture, with a thin client and swappable data and location tiers, kept the system adaptable: the data store and the provider source can each be replaced without touching engine, route, or interface code.

8.2 Recommendations

The immediate priority is security hardening of the API, since the system is otherwise feature-complete: restrict cross-origin access, add security headers, apply rate limiting to the unauthenticated endpoints, clamp request parameters, and bound the in-memory cache and query log. Next, implement the Firestore data store against the existing interface so the query log survives restarts, and wire live Maps credentials while keeping the key server-side. Clinical review of the knowledge base by a qualified practitioner should precede any real use, and the reviewer and date should be recorded. Beyond that, adding authentication would make history per-user, a continuous-integration pipeline would validate every change automatically, and multilingual and offline support would extend reach in the target context.

8.3 Closing Remarks

The system delivers a working, testable, and honest prototype: honest in that its limitations are documented as precisely as its capabilities, and that it tells every user, on every response, that it offers guidance rather than diagnosis. Acting on the recommendations above would carry it from a sound prototype to a deployable product.

References

Google Developers. (2024). *Maps Platform documentation*.

<https://developers.google.com/maps/documentation>SMARTPANTS PRESERVED 1725SMARTPANTS PRESERVED 1726

Semigran, H. L., Linder, J. A., Gidengil, C., & Mehrotra, A. (2015). Evaluation of symptom checkers for self-diagnosis and triage: Audit study. *BMJ*, 351, h3480.

<https://doi.org/10.1136/bmj.h3480>SMARTPANTS PRESERVED 1731SMARTPANTS PRESERVED 1732

World Health Organization. (2023). *Primary health care*. <https://www.who.int/health-topics/primary-health-care>SMARTPANTS PRESERVED 1737SMARTPANTS PRESERVED 1738

World Health Organization. (2022). *Global strategy on digital health 2020-2025*.

<https://www.who.int/publications/i/item/9789240020924>SMARTPANTS PRESERVED 1743SMARTPANTS PRESERVED 1744

Appendices

Appendix A: Repository

The complete source code, including the backend, the mobile client, the design figures, and the report sources, is available at github.com/AmyNjau/SWE3090XA-Project.

Appendix B: Running the System

The backend is started from the `backend/` directory with `npm run setup` followed by `npm start`, which serves the API on port 3000. The mobile client is run against it with the API base URL supplied at build time. Full instructions are held in `LAUNCH.md` in the repository.

Appendix C: Knowledge Base Coverage

The delivered knowledge base covers 15 conditions: malaria, common cold, influenza, COVID-19, migraine, gastroenteritis, hypertension, asthma, allergic rhinitis, pneumonia, urinary tract infection, tonsillitis, dengue fever, typhoid fever, and conjunctivitis. These are mapped across 9 specialist types through 35 distinct symptoms.